

Text modality model pipeline

Bilingual text is encoded by a fine-tuned XLM-RoBERTa classifier in a shared seven-emotion label space.

01
Text input

Chinese / English

?I am delighted!?

strip whitespace
reject empty input

02
Tokenization

SentencePiece

<s> I am delighted ...

input_ids / attention_mask
truncate to 128 tokens

03
Context encoding

XLM-RoBERTa encoder

Transformer block ?12

12 attention heads / hidden size 768

04
Sequence summary

<s> representation



context vector h (768-d)

05
Classification

Linear head

768 to 7 logits

softmax

class probabilities

06
Prediction

anger
disgust
fear
* joy
sadness
surprise
neutral
argmax + confidence

Inference

local model / evaluation mode / no gradients / CPU-GPU auto-selection / bilingual label space